

Milner's scheduler

Adrián Enríquez Ballester

December 9, 2021

We want to specify a simple scheduler for a set of n agents P_1, \dots, P_n . Each agent P_i performs a task repeatedly, and the scheduler is required to ensure that they begin the task in cyclic order starting with P_1 . The different task-performances need not exclude each other in time. For example, P_2 can begin before P_1 finishes, but the scheduler is required to ensure that each agent finishes one performance before it begins another.

We assume that P_i requests task initiation by an action a_i and signals completion by an action b_i . The scheduler can then be specified by requiring that:

- It must perform a_1, \dots, a_n cyclically, starting with a_1 .
- It must perform a_i and b_i alternately for each i .

However, a scheduler which imposes a fixed sequence, say $a_1b_1a_2b_2\dots$ is not good enough, the scheduler must allow any sequence of actions compatible with the two conditions above. For example, for $n = 2$, the sequences $a_1a_2b_1b_2a_1$ and $a_1b_1a_2a_1b_2b_1$ are compatible with the specification, but the sequences $a_1b_1a_1$ and $a_1b_1a_2a_1a_2$ are not.

1 Petri net model for two processes

The Figure 1 shows a diagram of the Petri net for the Milner's scheduler where $n = 2$. Each process has two places which represent its two possible states (i.e. idle or working), and two transitions a_i and b_i for the corresponding actions of starting and ending its task.

The scheduler state has been modeled with one place per process, which reflect the one that has to start next. The token changes from a place to the next one when the corresponding a_i transition is fired, allowing only the one which corresponds to the current state.

2 Properties

If we traverse the reachability graph of the net at the Figure 1, we reach 8 different markings. Every transition appears in the graph and we can reach again the initial marking from any other, so the net is live.




Figure 1: Petri net diagram for the Milner's scheduler of two processes

Also, as intended for properly modeling the states, every place is 1-bounded, thus the net itself is 1-bounded.

3 Generalization

The Petri net model can scale to n processes in the following way:

- For each $i \in \{1, \dots, n\}$, a process is represented with two places for its idle and working state, and two transitions a_i and b_i connected as in the example for $n = 2$. For the initial marking we can put a token in its idle place.
- For each $i \in \{1, \dots, n\}$, a place p_i for the scheduler sequence state. p_i should be a precondition of a_i and a postcondition of a_{i-1} (a_n when i is 1). For the initial marking we can put a token in a_1 .

4 Reachable markings growth

Each process can be either idle or working, and for each state of the scheduler sequence it is possible to reach any combination of the process states. For a Milner's scheduler of n agents, as there are n agents with 2 states each one, and

the scheduler sequence has n states, the amount of reachable markings of its Petri net (N, M_0) is as follows:

$$|[M_0]| = 2^n \cdot n$$